HACKATHON 2.0 — PROBLEM STATEMENTS | 39 Challenges | Points = Complexity Weight · Higher Points = Harder Problem = Greater Multiplier Reward

Select ONE problem from the list below. Each participant works individually . Higher base points = more ambitious challenge = greater reward multiplier.

#	Function	Project Name	Project Objective (What to Build)	Points (Weight)	AI Fit
1	Analytics	AI Business Review & Data Intelligence Platform	End-to-end meeting intelligence + data quality engine: auto-gathers WBR/MBR metrics, surfaces anomalies, generates pre-read decks, ingests AI note-takers, tracks action items, and learns stakeholder preferences over time.	72	
2	Customer Experience	AI Personal Concierge — Unified Post-Sales Bot	Unified post-sales AI bot (voice + chat) that handles all post-sales journeys, remembers context across conversations, proactively reaches out, and auto-resolves Salesforce cases with LLM+RAG.	78	
3	Customer Experience	Smart EDD — ML Delivery Date Prediction	ML model for accurate delivery date prediction — improve EDD accuracy from 45% → 80% using historical shipment data, logistics patterns, and real-time signals.	35	
4	Customer Experience	AI Detractor Predictor & Auto-Delight	Predict likely detractors from interaction signals and auto-trigger delight interventions — reduce detractors from 18 → 12 per 1000.	25	
5	Stores / CX / International	In-Store AI Assistant + Multilingual CMS	Unified in-store QR→WhatsApp AI concierge handling product queries, appointments, and post-visit follow-up — with multilingual layer for UAE/TH/SG/KSA.	55	
6	Stores	Store Ticket Auto-Resolution Agent	AI ticket classifier to reduce store tickets 850/day → 550/day with auto-resolution of task tickets and self-serve tools for common issues.	20	
7	B Glass	Smart Glass Agentic Companion on the app .	A voice agent that runs entirely on-device for the majority of customer turns, with a learned escalation router that hands off to a cloud LLM (Gemini / Claude) only when the on-device model lacks confidence or capability. The target architecture	95	
8	SCM	Tier 2/3 Delivery Speed — AI Cutoff Optimizer	ML-based courier cutoff logic for tier 2/3 cities — add intelligent optimization layer on existing dev for faster delivery speed.	22	
9	Merchandising	AI-Powered In-House Buying Platform	In-house forecasting and buying platform with ML signals — similarity index, Claire data, and customer signals for demand prediction.	60	
10	Merchandising / Catalog	AI Catalog Sanitation & Taste Agent	Multi-layered AI agent for catalog QC: image validation, classification audit, pricing sanity, attribute completeness, fashion score, and image authenticity — runs continuously.	65	
11	International / Technology	Multilingual Product & Content CMS	CMS with LLM-powered auto-translation for catalog/content across Arabic, Thai, etc. — with maker-checker approval UI and translation coverage tracking.	40	
12	Finance	AP Invoice Auto-Validation Agent	AI agent that OCRs vendor invoices via Claude Vision, matches against PO & GRN in D365/FEBI, flags discrepancies, and drafts posting entries.	45	
13	Marketing	AI Influencer Discovery & ROI Predictor	Tool that ranks influencers by audience affinity, predicts campaign ROI, and scores video quality — cutting discovery TAT from 2 days → 3 hours.	55	
14	Marketing	Creative Generation & Social Listening Suite	Auto-generate & A/B test Google/Meta creatives + continuous social listening agent scanning X/Instagram/Reddit for brand mentions and sentiment.	58	
15	Online	AI Shopping Experience — Claire + Face Personalization	Unified AI shopping assistant: frame quiz, face-based recommendations, lens guidance, face-scan PLP ordering, and photo-based size-confidence widget.	82	
16	SCM	NexS Order Prioritization Engine	Order-prioritization logic inside NexS so delayed/high-EDD-risk orders process first at plant & LF — feeding NDD 99% and SV consistency goals.	38	
17	SCM	Tail-Tray RFID Locator + Anomaly Agent	Dashboard + agent cross-referencing RFID scans with order IDs, finding missing/tail trays in real time, and alerting ops for NDD & SV 99%+.	25	
18	Technology	Observability Co-Pilot for SREs	LLM that ingests Grafana/Datadog/OpenTelemetry traces and proposes root-cause + fix PR for incidents — cutting MTTR further.	62	
19	Technology	FraudLens Threshold Auto-Tuner	Auto-calibration agent that learns optimal FraudLens thresholds from recent data + human-in-loop reviews — unblocks revenue-loss prevention at scale.	42	
20	Merchandising	Trend-to-SKU Translator	Watches fashion signals (Vogue, Pinterest, TikTok, K-drama) and translates emerging aesthetics into vendor-ready SKU spec sheets with reference images and price bands.	65	
21	Online	Face-Scan Twin → Try-On Video Agent	Given face scan + frame selection, generate a short video of customer wearing frame in 15 different lighting/scene contexts in <30 seconds, shareable via WhatsApp.	30	
22	Online	Lens Recommendation Engine	AI-powered lens recommendation based on Rx complexity, lifestyle questionnaire, budget, and usage patterns — replacing manual optician guidance for lens selection.	42	
23	Post-Sale / CX	Return Reduction Diagnostic Agent	Agent that conducts targeted questions on return initiation, classifies root cause, offers tailored save-the-sale (exchange/lens redo/optom callback), and feeds data upstream.	48	
24	Offline / Stores	Self Checkout Kiosk	Self-checkout experience for stores — customer scans products, pays via UPI/card, walks out. Reduces checkout queue time and frees store associates for consultative selling.	88	
25	Offline / Stores	Store Recommendation Engine (Beat Current)	New in-store recommendation engine backtested against current one — face geometry + style preferences + purchase history. Must demonstrably beat existing system.	70	
26	Offline / Stores	Store Performance Diagnostic Agent	Agent identifying conversion change drivers per store, generating actionable playbooks for degrowing stores with automated store health reports for ZMs.	35	

27	SCM	Remote Edging Kiosk Prototype	Store captures PD/Rx/lens, sends to central lab, returns fitted glasses overnight — demo with 1 pilot store for LF-less city delivery.	45	
28	People	Employee Grievance & Sentiment Agent	Agent to intake, classify, route, and track employee grievances + real-time sentiment scoring from pulse surveys/gmeet with HRBP escalation.	28	
29	Legal	Compliance Training Auto-Generator	Ingests latest regulations, generates role-specific training modules, creates quizzes, and tracks completion across ~6000 employees — zero-touch annual compliance cycle.	22	
30	Finance	D365 PR-to-PO Co-Pilot + Vendor Onboarding	Agent that autocompletes vendor/SKU from PO history, predicts approval bottlenecks + conversational vendor KYC agent collecting docs and auto-creating vendor master.	38	
31	Moonshot	Eye Test on App	Clinically-credible phone eye screening (visual acuity, near vision, contrast sensitivity, refractive error) with LLM report synthesis and auto-booking of nearest optom slot.	92	
32	Moonshot	Autonomous Store Ops Replay	Ingest 48 hrs of multi-camera CCTV + audio; produce a 3-minute narrated ZM video summary with key incidents, customer flow, staff performance highlights, and actionable callouts.	90	
33	Builder Productivity	Ship-a-Feature Agent	PM writes 3-line feature spec → agent produces PR across frontend + backend + tests + feature flag config + staging preview URL within 30 minutes.	85	
34	Builder Productivity	Bug-to-PR Agent	Auto-converts Jira/Linear bug reports into PRs with fix + tests + changelog entry. Triages severity, identifies root-cause file, generates fix, runs tests, submits for review.	58	
35	Offline / Stores	Store Staff Allocation & Conversion Agent	Ingests real-time signals (walk-in volume, staff skills, queue lengths) to dynamically recommend who should do what, when — surfaces via SM app/WhatsApp.	62	
36	Post-Sale / CX	Comms Channel Deviation Scanner	Agent scanning emails/Slack/Teams for deviations from expected patterns (missed deadlines, compliance red flags, sentiment shifts) with escalation and context.	55	
37	Technology / Pricing	Pricing Rules Sandbox & Safe-Deploy Agent	Build an AI-powered sandbox environment that intercepts any new pricing rule before it hits production and automatically runs a full-impact simulation: (1) tests the rule against the last 90 days of real order data to forecast revenue delta, margin impact, and edge-case anomalies; (2) flags conflicting rules (e.g. a new discount that stacks unintentionally with an existing promo), price-floor breaches, and category-level inconsistencies; (3)	60	
38	Builder Productivity	PageDrop	Build a tool/platform where a non-engineer can vibe-code a product page (using any AI-assisted method) and publish it directly into the Lenskart app as a native-feeling page — not a janky webview, not an iframe hack. The tool itself must handle all pre-production sanity (responsive rendering, performance, design-system compliance, broken assets, API contract checks) so that the builder doesn't need to become a QA engineer to ship	70	
39	Moonshot	Vibe Code to Production	ple across Product, Ops, Merchandising, and CX regularly prototype useful tools and features using AI-assisted coding. But these prototypes die in staging because deploying to production requires engineering handoff, infra provisioning, security review, and CI/CD setup — none of which the builder controls.	95	

POINTS GUIDE: 85–95 = Very Hard (Moonshot) 65–84 = Hard 50–64 = Medium-Hard 35–49 = Medium 18–34 = Focused / Starter | Higher Points = bigger multiplier reward when executed well. | AI Fit = AI is the whole product = AI powers the core logic

POINTS DISTRIBUTION — Pick your complexity level		
85–95 pts	Very Hard	#29 Self-Checkout (88) #36 Eye Test on App (92) #38 Ship-a-Feature (85) #44 Vibe Code to Prod (95)
65–84 pts	Hard	#1 AI Biz Review (72) #12 Store Zone CCTV (80) #19 AI Shopping (82) #20 Multi-Modal (75) #25 Trend-SKU (65) #30 Store Rec Engine (70) #43 PageDrop (70)
50–64 pts	Medium-Hard	#5 In-Store Asst (55) #7 2D Try-On (50) #10 Buying Platform (60) #11 Catalog Agent (65) #17 Influencer (55) #18 Social (58) #23 Observability (62) #37 Store Replay (55) #39 Bug-to-PR (58) #40 Staff Alloc (62) #41 Comms (55) #42 Pricing Sandbox (60)
35–49 pts	Medium	#3 Smart EDD (35) #14 Multilingual CMS (40) #15 AP Invoice (45) #21 NexS Order (38) #24 FraudLens (42) #27 Lens Rec (42) #28 Return Diag (48) #31 Store Perf (35) #32 Remote Edging (45) #35 D365 Co-Pilot (38)
18–34 pts	Focused	#4 Detractor (25) #6 Store Tickets (20) #8 Omni Dash (18) #9 Cutoff Opt (22) #13 360 Feedback (30) #16 Compliance Mon (28) #22 RFID (25) #26 Try-On Video (30) #33 Grievance (28) #34 Training Gen (22)

HACKATHON 2.0 — DETAILED PROGRAM SCHEDULE | 23rd & 24th April 2026			
Time Slot	**Session / Activity**	**Owner / Speaker**	**Notes for Participants**
DAY 1 — Thursday 23rd APRIL 2026 | KICK-OFF & BUILD BEGINS			
9:30 AM – 10:00 AM	**OFFICIAL KICK-OFF ADDRESS**	**Ramneek & Natraj (Leadership)**	Mandatory for all participants. Leaders will set context, announce prizes, and energise the room.
10:00 AM Onwards	BUILD PHASE BEGINS — Individual Development Sprint	**Participants**	Build starts now! You can begin building from 9 AM if you want — the check-in, kick-off, and briefing sessions are strongly recommended but you're free to dive in whenever you're ready. Mentors are on the floor all day — grab any mentor anytime.
DAY 2 — Friday, 24th APRIL 2026 | FINAL SPRINT & SUBMISSION			
2:00 PM – 4:30 PM	Final Build & Documentation Sprint	**All Participants**	Polish your agent. Write your 1-pager approach note. Finalise your GitHub repo. Record your 3-minute demo video.
4:30 PM – 5:00 PM	**SUBMISSION WINDOW — FINAL DEADLINE: 5:00 PM SHARP**	**All Participants**	HARD DEADLINE. Late submissions not accepted. Double-check all three components are uploaded.
5:00 PM – 6:00 PM	Judge Orientation + Scoring Calibration	**Judges + Program Leads**	Judges review submissions.
6:00 PM – 7:30 PM	**LIVE DEMOS — Top Shortlisted Submissions (3 min demo + 2 min Q&A each)**	**Judges + Shortlisted Participants**	If shortlisted, you'll be notified at 5:30 PM. 3-minute live demo + 2 minutes of judge questions.
7:30 PM – 8:00 PM	**WINNERS ANNOUNCEMENT + CLOSING CEREMONY + CERTIFICATES**	**Leadership (Ramneek & Natraj)**	Prizes, certificates, photo opportunity, and wrap-up. Everyone welcome.

EXTERNAL SPEAKER	There will be sessions with external speakers throughout Thursday (23rd April) and Friday (24th April)
Food & Beverages	Breakfast (24th Morning) , Lunch (23rd & 24th Afternoon) & Dinner (23rd Night) will be served to Participants in office . Hot & Cold Beverages will be served throughout 2 days .

YOUR SCORING GUIDE — HACKATHON 2.0

FINAL SCORE = Base Points × Core Multipliers × Bonus Multipliers | Read this before you build

HOW SCORING WORKS: Every problem has a Base Point value (18–92 pts). Your final score = Base Points × all applicable multipliers. Higher base = harder problem = bigger reward when you execute well. Multipliers stack — a great demo on a hard problem can produce a score 10x the base.

PART 1 — BASE CRITERIA (3 criteria, 10 pts each, max 30 pts)

① PROOF OF WORK "Does it actually run?"	**10 / Exceptional** **Live end-to-end demo on real or realistic data. Agent handles at least one unexpected input without intervention. Zero crashes.**	**In your VIDEO** **Show a LIVE run — raw input goes in, agent reasons on screen, output comes out. No cuts, no pre-recorded clips. Give it one input it hasn't seen before.**
	7–9 / Good Works end-to-end. Minor rough edges (one retry, basic UI) but the core agent completes its full task.	In your CODE README must have working setup instructions (clone → configure → run in ≤ 5 commands). Judges will try to run it.
	4–6 / Adequate Core happy path runs but breaks on edge inputs, or demo data is obviously hardcoded/mocked.	In your NOTE Describe the agent's workflow in 3–5 bullet points. Does what you wrote match what the video shows? Inconsistency costs points.
	1–3 / Weak Only a fragment of the agent runs, or the demo is a pre-recorded clip of an earlier run.	
	0 / No demo Slides, mockups, or concept decks only. No working agent shown. Automatic zero — no exceptions.	
② REAL IMPACT "Does it move a number?"	**10 / Exceptional** **Quantified impact evidenced in both the note AND the demo: 'Invoice TAT: 3 days → 2 hours (85% faster)' — shown as a before/after moment on screen.**	**In your NOTE** **Write exactly ONE sentence in this format: "Before: [metric]. After: [metric]. Delta: [number + unit]."** **Example: Before: 3-day manual review. After: 20-min automated check. Delta: 85% faster.**
	7–9 / Good Named metric with a directional number (e.g. '~60% fewer manual steps'). Shown in the demo even if not perfectly precise.	In your VIDEO Show the BEFORE state first (the broken/slow thing), then show the agent fixing it. Both states must appear on screen — don't just describe the improvement.
	4–6 / Adequate Real pain point described but impact is qualitative only — 'saves time', 'reduces errors' — no number anywhere in the submission.	The unit rule Use a real unit: ₹, %, hours, days, tickets, errors, SKUs, clicks. If your submission has zero units anywhere, evaluators cap the score at 5.
	1–3 / Weak Problem is mentioned but the agent doesn't visibly solve it in the demo.	
	0 / No value No articulation of who benefits or by how much.	
③ AI IS THE ENGINE "Is AI doing the real work?"	**10 / Exceptional** **Multi-step reasoning shown: agent takes input → decides what to do → calls a tool or data source → synthesises output → handles a different input differently. Couldn't exist as traditional if/else software.**	**In your VIDEO** **Show TWO contrasting inputs back-to-back that produce meaningfully different outputs. This is the clearest proof the AI is reasoning, not pattern-matching.**
	7–9 / Good LLM or ML model is clearly the decision-maker in the core workflow. Replacing it with rules would break the product.	In your CODE Include your system prompt or agent config in the README (even a 5-line snippet). A thoughtful, domain-specific system prompt is strong evidence of AI centrality.
	4–6 / Adequate AI is used but could be replaced by a lookup table or simple rule. It adds convenience, not intelligence.	The test Ask yourself: if I replaced the LLM with a lookup table and 20 if/else rules, would it still work? If yes → your AI score will be low. Make the AI indispensable.
	1–3 / Weak AI was used to write the code (Cursor, Claude Code) but the product itself has no LLM/ML in the runtime path.	
	0 / No AI No AI or ML in the final product at all.	

PART 2 — SCORE BANDS (before multipliers)

27–30	SHIP IT	All three dimensions excellent. Present at All-Hands. Immediate productionisation path.
21–26	STRONG BUILD	Two dimensions excellent, one good. Assign a sponsor, 2-week sprint to go live.
14–20	PROMISING	Works and has a valid idea. Missing quantified impact or AI depth. Add to backlog with a clear owner.
7–13	REBUILD	Core concept valid but execution gaps are significant. Good learning — not ready to ship.
0– 6	INCOMPLETE	No working demo or no discernible impact. Participation noted but not ranked.

PRO TIPS FROM THE JUDGES

1. Record your demo video at the END of Day 1 — so Day 2 is for polishing, not panicking.
2. Show the BEFORE state in your demo (the broken thing) before showing the fix. Without a before, there's no impact.
3. Paste your system prompt in the README. It takes 30 seconds and is the single strongest signal for the AI Engine criterion.
4. Say the number out loud in the video. Don't leave it buried in the note.

PART 3 — MULTIPLIERS (these stack on top of your base criteria score)

Every applicable multiplier multiplies your score further. A 2.0× multiplier doubles your score. Multiple multipliers stack — the best submissions can reach 10× their base score.

MULTIPLIER / CONDITION	WHAT IT MEANS — HOW TO EARN IT	FACTOR
	CORE MULTIPLIERS (all six apply to every submission)	
① DEPLOYABILITY	How quickly can this go live in production? 2.0× → Deploy right after hackathon: working end-to-end, connected to real systems, can switch on immediately 1.5× → Deployable within 1 week: core flow works, needs config / auth / edge-case handling 1.2× → Needs 2–3 weeks of hardening: proves concept but needs integration and testing 0.5× → Month+ to production: significant gaps — treat as R&D / proof of concept	**0.5× – 2.0×**
② USE CASE COVERAGE	What % of the original problem goal does your output cover? 2.0× → 100%+ of original scope: all stated KPIs addressed, plus bonus capabilities beyond scope 1.5× → 70–100% of scope: core use cases work, maybe missing edge cases or secondary flows 1.0× → 50–70% of scope: primary happy path works but significant gaps remain 0.5× → Sub-50%: narrow slice only — penalises over-scoping without delivering	**0.5× – 2.0×**
③ AI NATIVENESS	How deeply is AI / LLM / agentic tech embedded at the core? 1.8× → AI IS the product — agentic, autonomous: multi-step reasoning, RAG, tool-use, feedback loops 1.4× → AI powers the core logic: ML model or LLM does the heavy lifting within a traditional shell 1.0× → AI assists (copilot, suggestions): AI adds value but solution would partially work without it 0.5× → AI only used in building: used Cursor / Claude Code to build faster, but product is traditional	**0.5× – 1.8×**
④ ORG LEVERAGE	Can other teams / functions reuse this with minimal effort? 1.5× → Platform usable by 3+ teams as-is: e.g., a review engine any function plugs into 1.2× → Reusable with config for 1–2 teams: core architecture is extensible, needs some customisation 0.8× → Single-team use only: solves one problem well but doesn't scale to other teams	**0.8× – 1.5×**
⑤ BUSINESS IMPACT	How directly does this move revenue, cost, or customer experience? 2.0× → Direct P&L or NPS impact, quantified: team shows '₹X/month saved' or 'Y% fewer tickets' with real data 1.5× → Strong leading indicator: clear operational improvement — faster TAT, fewer errors — P&L link indirect 1.0× → Internal productivity improvement: makes teams faster but doesn't directly touch customer/financial metrics 0.3× → Unclear business value: can't articulate who benefits and by how much	**0.3× – 2.0×**

⑩ DEMO QUALITY	Did you ship something that feels real? 1.5× → Polished, production-feel UX: clean UI, handles errors gracefully, no 'let me restart' moments 1.0× → Functional but rough prototype: core flow works end-to-end but basic UI, some manual steps 0× → Slides / video only — no working product: no working demo = ZERO SCORE regardless of idea quality	0× – 1.5×

BONUS MULTIPLIERS (optional — earn any that apply, they all stack further)		
AI Adoption Transfer	You taught ≥ 2 non-technical colleagues to use your tool or modify prompts during the hackathon.	×1.3
Speed Record	Working demo completed and shown in under 6 hours from build start (Day 1 afternoon checkpoint).	×1.3
Cross-Function Build	You actively collaborated with someone from a different function (e.g., Tech + Finance, Marketing + SCM).	×1.2
Documentation	You published a runbook, prompt guide, or setup doc so anyone in the company can replicate your work.	×1.1
Live Data	Your demo runs on real production data — not mocked, synthetic, or hardcoded sample data.	×1.2

SCORE TIERS (after all multipliers applied)		
LEGENDARY	Exceptional on every dimension — immediate adoption, transformative impact, org-wide reuse. Present at All-Hands.	2,000+
OUTSTANDING	Strong across the board — ready to deploy, high AI nativeness, clear business value. Fast-track to production.	500–2,000
SOLID	Good prototype with production gaps — needs a follow-up sprint but promising. Add to backlog with a sponsor.	100–500
NEEDS WORK	Significant gaps — good learning exercise but not ready to ship. Revisit fundamentals.	< 100

WORKED EXAMPLE — AP Invoice Auto-Validation Agent (Problem #15, Finance | Base: 45 pts)

SCENARIO: Priya (Finance) built an agent that OCRs vendor invoices, matches them against PO & GRN in D365, flags mismatches, and drafts posting entries. She showed it live in a demo on realistic invoice PDFs. Her 1-pager states: 'Before: 3-day manual validation. After: 20-min automated check. Delta: 85% faster.' Her GitHub README includes her system prompt. The agent is fully working but only covers the India D365 instance — UAE/KSA are not included. She built it with her manager who is from the SCM team (cross-function).

STEP	SCORE / MULTIPLIER	REASONING — WHY THIS SCORE?
BASE CRITERIA ① Proof of Work	8 / 10	Live demo runs end-to-end on realistic invoice PDFs. Agent correctly flags a mismatch and drafts the posting entry on screen. Minor rough edge: one edge-case invoice caused a timeout (she restarted). README has working setup in 4 commands. Deduction of 2 pts for the timeout and pre-production data.
BASE CRITERIA ② Real Impact	9 / 10	'Before: 3-day manual validation. After: 20-min automated check. Delta: 85% faster.' This appears in both the 1-pager AND is shown live — the demo opens with the slow manual flow, then cuts to the agent completing the same task. All units present. Loses 1 pt because the ₹ savings figure is estimated, not from actual production run.
BASE CRITERIA ③ AI Engine	8 / 10	LLM is in the critical path: it reasons about mismatches between OCR'd invoice fields and D365 data, generates a natural-language flag ("Vendor quoted ₹12,400 but PO says ₹11,800 — likely GST discrepancy"), and drafts the posting entry. Two contrasting invoices shown: one clean pass, one flagged with explanation. System prompt in README is domain-specific and detailed. Minus 2 pts: no tool-use or RAG, single LLM call.
CRITERIA TOTAL	25 / 30	8 + 9 + 8 = 25 pts. Strong across all three dimensions. Now the multipliers amplify this.
BASE POINTS × CRITERIA (Problem #15 = 45 pts base)	45 pts	The base point value of the problem (set by problem difficulty) is separate from the criteria score. Final score = Base Points × all multipliers. The 25/30 criteria score feeds into the multipliers below — it establishes which multiplier tier you land in for each dimension.
① DEPLOYABILITY	× 1.2	Core flow works, connected to real D365 test instance, but only India. UAE/KSA not covered. 2–3 weeks of hardening needed before full production switch-on. → Not × 1.5 (would need all geographies working) but better than × 0.5 (this is clearly productisable).
② USE CASE COVERAGE	× 1.5	Covers OCR → match → flag → draft posting for India. That's 70–90% of the original problem scope. Missing: UAE/KSA, auto-approval workflow, and the vendor onboarding part of the brief. → Lands at × 1.5 (70–100% of scope), not the full × 2.0.
③ AI NATIVENESS	× 1.4	LLM is doing the core reasoning and generation, not just formatting text. A rule-based system could not produce the natural-language mismatch explanation. → × 1.4 (AI powers the core logic). Just below × 1.8 because there's no tool-use, RAG, or multi-step reasoning loop — it's a single inference call with a strong prompt.
④ ORG LEVERAGE	× 1.2	The OCR + D365 matching architecture could be adapted for other finance flows (PO validation, vendor onboarding) with 1–2 weeks of customisation. One team could reuse it. → × 1.2 (reusable with config for 1–2 teams). Not × 1.5 because it's not a plug-and-play platform yet.
⑤ BUSINESS IMPACT	× 1.5	Clear operational improvement — 85% faster TAT is a strong leading indicator. The P&L link is obvious (fewer FTE hours, fewer posting errors, faster book closure) even if not yet quantified in ₹. → × 1.5 (strong leading indicator). Would need actual ₹ saved from a production run to earn × 2.0.
⑩ DEMO QUALITY	× 1.0	Functional and complete end-to-end but UI is basic (terminal output + a simple form). One restart happened during the demo. The before/after structure is clear. → × 1.0 (functional prototype). Would need clean UI and zero restarts for × 1.5.
BONUS: Cross-Function Build	× 1.2	Built with her SCM manager — Tech (Priya, Finance) + SCM. Confirmed cross-function collaboration. Bonus applies cleanly.
FINAL SCORE	81	45 × 1.2 × 1.5 × 1.4 × 1.2 × 1.5 × 1.0 × 1.2 = 81 pts → SOLID VERDICT: Strong prototype with real impact. Add to Finance backlog with Nishant as sponsor. Recommended next step: extend to UAE/KSA D365, upgrade UI, quantify ₹ savings from a 2-week pilot. With those changes this hits OUTSTANDING easily.

KEY LESSON FROM THIS EXAMPLE: Priya's agent scored SOLID (81 pts) not because any single thing was wrong — her criteria scores were excellent (25/30). The score reflects that she only covered India, had a basic UI, and didn't quantify ₹ savings from production. Each of those is fixable in a 2-week sprint. The multipliers reward completeness and production-readiness, not just proof-of-concept quality.

SUBMISSION GUIDELINES — HACKATHON 2.0

HARD DEADLINE: 5:00 PM, Friday 24th April 2025

COMPONENT 1 — 1-PAGER APPROACH NOTE		
What it is	A single-page document (PDF or Google Doc) summarising your agent, the problem it solves, and how it works.	
Problem Statement	Which of the 41 problems did you pick? What specific pain point does your agent address? What is the baseline today (metric)?	
Your Solution	What does your agent do? Describe the workflow in 3–5 bullet points. What AI/LLM technique did you use (RAG, tool-use, ReAct loop)?	
Tech Stack	List the tools and frameworks: Open Claw, AWS Bedrock model(s), SageMaker, any APIs, data sources used.	
Business Impact	What metric does this move? By how much? E.g., 'Reduces invoice validation TAT from 3 days to 2 hours' — quantify wherever possible.	
Assumptions & Limitations	What did you assume? What would need to change for this to go to production?	
Format	Max 1 page. Can be a Google Doc, Notion page, or PDF. Use the template provided in the submission portal.	

COMPONENT 2 — CODE SUBMISSION (GITHUB LINK)		
What to submit	A GitHub (or GitLab) repository link. The repo must be accessible to judges (public OR invite judges@hackathon.internal as collaborator).	
README.md	Your repo MUST have a README with: (a) Setup instructions — how to run the agent locally, (b) Environment variables needed (use .env.example, never hardcode secrets), (c) Architecture diagram or flowchart, (d) Sample inputs and expected outputs.	
Code Quality	Clean, commented code. No secrets in repo. Use .gitignore. Judges will look at commit history — show your work progressively.	
Dependencies	Include requirements.txt (Python) or package.json (Node). The agent should be runnable with 3 commands or fewer.	
Branching	Main/master branch = final submission. Any experimental branches are fine but judges evaluate main only.	
Data	If you used real production data, anonymise it. Synthetic/mock data is perfectly fine and preferred for demos.	

COMPONENT 3 — 3-MINUTE DEMO VIDEO		
What it is	A screen-recorded video (with voiceover) showing your agent in action. Maximum 3 minutes. No slides — show the actual product working.	
What to show	1. The problem (10 seconds) — what was broken or slow before? 2. Your agent in action (2 minutes) — live demo of the core workflow. Show real inputs and real outputs. 3. The impact (30 seconds) — what's better now? Show the metric.	
Recording tools	Loom (recommended), OBS Studio, QuickTime, or any screen recorder. Must include audio narration.	
File format	MP4 preferred. Max 500 MB. Upload to Google Drive / Loom and submit the shareable link (ensure judges can view without account).	
Language	English preferred. Hindi acceptable. Subtitles a bonus.	
No 'coming soon'	Do NOT show slides of planned features. Show what WORKS. A simple working agent beats an ambitious demo of mock data.	

PRE-SUBMISSION CHECKLIST — Tick all before hitting Submit
☐ 1-Pager Approach Note is written, formatted, and exported as PDF / shareable Google Doc link
☐ GitHub repo is accessible to judges (public or collaborator invite sent)
☐ README.md is complete: setup instructions, architecture, sample I/O, no secrets in code
☐ 3-minute demo video is recorded, uploaded to Google Drive/Loom, and share link works
☐ All three submission links pasted into the google drive link before 5:00 PM on 24th April